

Synthesis Project — Questions & Answers

iDISC Translator Recommendation System

A. Data Sufficiency & Cold-Start

Q1

Is it normal for you as a company to have a high rotation of translators, meaning they only get to do a few tasks for iDISC? Were they temporary workers, or was it a special condition? Do you consider them for the assignment of daily tasks?

Answer

Yes, this situation is normal and structural in our operating model. We typically have a long tail of translators with low activity. This is not mainly due to temporary contracts, but to onboarding tests, niche language pairs, irregular demand, or availability constraints. Many translators join the system, while only a smaller group becomes regular contributors.

From an operational standpoint, daily task assignment prioritises translators with sufficient history, reliability, and availability. Low-activity profiles are used mainly in specific or exploratory cases, not as default options. In short, high rotation at the tail is expected, reflects business dynamics, and these profiles are not weighted equally in routine assignment decisions.

Q2

If we must include profiles of translators that have a low history of tasks completed, we could face a huge problem when trying to confidently recommend a translator for a given task. We are considering including a threshold of a minimum of tasks completed to take part in the training data. Is this approach valid?

Answer

Your concern is operationally valid and well framed from a client perspective. Profiles with very limited task history simply do not provide enough evidence to support confident assignment decisions. This is not a technical issue, but a natural consequence of how work history accumulates in operations.

Introducing a minimum-activity threshold is therefore a legitimate way to protect decision quality. For your analysis, the focus should be on how such a rule balances reliability and coverage in daily operations, rather than on the exact threshold value. Think in terms of decision confidence, operational risk, and how low-information profiles are typically managed.

B. Outliers & Distribution Anomalies

Q1

The top translator has completed 65,209 tasks while the average is 1,015. Are these high-volume translators a special category (e.g., agencies, bots, power users)?

Answer

From an operations perspective, extremely high task volumes do not necessarily mean 'more productive translators'. In many organisations, these profiles may correspond to full-time reviewers, checkers, or

validation roles that sit late in the workflow and touch a large number of files before delivery to the client.

So again, the key is not to assume a single explanation. For your analysis, you should consider: whether these profiles perform different operational functions (e.g. review, QA, final checks), whether they should be compared directly with standard translators, and how mixed roles impact averages, benchmarks, and assignment logic.

Q2

The mean total hours per translator is 1,752 but the median is only 133. What explains this extreme skew? How reliable is the HOURS column as a workload proxy?

Answer

This skew is consistent with operations where most volume is handled by a small, stable core, while specialised translators contribute fewer hours for specific needs.

HOURS remains the primary workload proxy, but it captures organisational choices as much as capacity. Part of the concentration is also explained by human scheduling bias: relying on familiar profiles is often operationally simpler than managing a wider pool, even if that limits diversification.

Q3

99 translator rates exceed the maximum client selling price (€50). Are these valid rates, legacy data, or errors?

Answer

(No comment provided in the document for this question.)

C. Language Pair Integrity

Q1

One translator covers 58 language pairs. Is this realistic, or does it indicate a data quality issue (e.g., an agency account)?

Answer

Coverage of 58 language pairs usually indicates an aggregated or non-standard account, such as an agency profile or a different process, like technical QA or a sign off review.

E. Task Type & Quality Labels

Q1

What is the full taxonomy of TASK_TYPE values? Are all task types expected to be automated, or only a subset?

Answer

Check file '260223_GraulAI_FieldExplanation.docx' for full list and description of tasks. Right now not expected to automate any of them.

Q2

How is the quality score defined and measured? Is it a numerical scale, a categorical rating, or a derived composite?

Answer

Qualitative analysis by reviewer based on specific guidelines. Ratings 1 to 10.

F. Temporal Validity

Q1

Are there translators who were active historically but are no longer available? How should inactive translators be flagged?

Answer

There is a boolean field in another table indicating whether a translator is available or not. This information cannot be provided at the moment.